

SCGC Assignment

by

ALEX APOSTOLESCU
aapostolescu

O carte citită e un grătar pierdut.
Jador

1 Task 0 - Cluster setup

Task 0 is accomplished using `./deploy-cluster.sh` which uses `kind` to create a node. I also included dependencies (`helm` and `slowhttptest`) for other tasks in this script to ensure the work environment is suitable.

```
student@lab-kubernetes:~/scgc$ ./deploy-cluster.sh
Helm v3.12.0 is already latest
Creating cluster "kind" ...
  ✓ Ensuring node image (kindest/node:v1.23.4) 🖼️
  ✓ Preparing nodes 📦
  ✓ Writing configuration 📄
  ✓ Starting control-plane 🚦
  ✓ Installing CNI 🖱️
  ✓ Installing StorageClass 💾
Set kubectl context to "kind-kind"
You can now use your cluster with:

kubectl cluster-info --context kind-kind

Have a nice day! 🍀
node/kind-control-plane condition met
NAMESPACE      NAME                                                    READY   STATUS    RESTARTS   AGE
kube-system     coredns-64897985d-2j9bh                                0/1     Pending   0           14s
kube-system     coredns-64897985d-86twc                                0/1     Pending   0           14s
kube-system     etcd-kind-control-plane                               1/1     Running   0           31s
kube-system     kindnet-4v5kw                                           1/1     Running   0           14s
kube-system     kube-apiserver-kind-control-plane                     1/1     Running   0           31s
kube-system     kube-controller-manager-kind-control-plane            1/1     Running   0           31s
kube-system     kube-proxy-vqgv8                                        1/1     Running   0           14s
kube-system     kube-scheduler-kind-control-plane                     1/1     Running   0           31s
local-path-storage  local-path-provisioner-5ddd94ff66-ftf7n              0/1     Pending   0           14s
student@lab-kubernetes:~/scgc$
```

After deploying the cluster, the script will wait up to 5 minutes for the nodes to be Ready.

2 Task 1 - Nginx

The base nginx is deployed using `./nginx/deploy-nginx.sh` which creates the configmaps for the main route and the metrics route, then applies the deployment, and starts a service that makes nginx available to the host.

```
student@lab-kubernetes:~/scgc/nginx$ ./deploy-nginx.sh
configmap/nginx-html-config created
configmap/nginx-metrics-config created
deployment.apps/nginx created
service/nginx created
pod/nginx-8ddb6569-hndt9 condition met

Test:
<html><body>Not everybody can be bombardier!</body></html>
Active connections: 1
server accepts handled requests
 2 2 2
Reading: 0 Writing: 1 Waiting: 0
student@lab-kubernetes:~/scgc/nginx$
```

Afterward, the script waits for the pods to be up and runs tests for both routes by running `curl` on `http://172.18.0.2:30080` and `http://172.18.0.2:30088/metrics`, where 172.18.0.2 is the node IP.

3 Task 2 - Promexporter

Deploying `promexporter` follows the same path: run the deployment, expose it via a service, wait for it to be online, and test it.

```
student@lab-kubernetes:~/scgc/promexporter$ ./deploy-promexporter.sh
deployment.apps/promexporter created
service/promexporter created
pod/promexporter-654d546d5d-8b976 condition met

Test:
<!DOCTYPE html>
<title>NGINX Exporter</title>
<h1>NGINX Exporter</h1>
<p><a href="/metrics">Metrics</a></p>student@lab-kubernetes:~/scgc/promexporter$
```

Testing is realized via curl on `http://172.18.0.2:30113`.

4 Task 3 - Prometheus

Prometheus is deployed in a namespace of its own, **monitoring** that will be created if it is missing.

After ensuring the namespace is present, **helm** is used to deploy **Prometheus** to this namespace and uses **values.yaml** to configure it. The only required configuration for our task are the scrape interval, the metrics path, and the target.

```
student@lab-kubernetes:~/scgc/prometheus$ ./deploy-prometheus.sh
registry.k8s.io/kube-state-metrics/kube-state-metrics:v2.8.0
namespace/monitoring created
"prometheus-community" already exists with the same configuration, skipping
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "prometheus-community" chart repository
...Successfully got an update from the "bitnami" chart repository
Update Complete. 🎉Happy Helming!🎉
NAME: prometheus
LAST DEPLOYED: Fri May 12 13:53:32 2023
NAMESPACE: monitoring
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
The Prometheus server can be accessed via port 80 on the following DNS name from within your cluster:
prometheus-server.monitoring.svc.cluster.local

Get the Prometheus server URL by running these commands in the same shell:
export POD_NAME=$(kubectl get pods --namespace monitoring -l "app=prometheus,component=server" -o jsonpath="{.items[0].metadata.name}")
kubectl --namespace monitoring port-forward $POD_NAME 9090

The Prometheus alertmanager can be accessed via port on the following DNS name from within your cluster:
prometheus-kls(<nil>).monitoring.svc.cluster.local


Get the Alertmanager URL by running these commands in the same shell:
export POD_NAME=$(kubectl get pods --namespace monitoring -l "app=prometheus,component=" -o jsonpath="{.items[0].metadata.name}")
kubectl --namespace monitoring port-forward $POD_NAME 9093
#####
##### WARNING: Pod Security Policy has been disabled by default since #####
##### it deprecated after k8s 1.25+, use #####
##### (index.Values "prometheus-node-exporter" "rbac" #####
##### "pspEnabled") with (index.Values #####
##### "prometheus-node-exporter" "rbac" "pspAnnotations") #####
##### in case you still need it. #####
#####

The Prometheus PushGateway can be accessed via port 9091 on the following DNS name from within your cluster:
prometheus-prometheus-pushgateway.monitoring.svc.cluster.local

Get the PushGateway URL by running these commands in the same shell:
export POD_NAME=$(kubectl get pods --namespace monitoring -l "app=prometheus-pushgateway,component=pushgateway" -o jsonpath="{.items[0].metadata.name}")
kubectl --namespace monitoring port-forward $POD_NAME 9091

For more information on running Prometheus, visit:
https://prometheus.io/
pod/prometheus-server-6cbbd55d6-tlph condition met
student@lab-kubernetes:~/scgc/prometheus$ Forwarding from 0.0.0.0:9090 -> 9090
```

Functionality is checked via the **Prometheus** UI, running queries on **nginx** statistics. This implies that for **Prometheus** to work correctly, the first 2 tasks should also be up and running.



5 Task 4 - Grafana

Grafana is somewhat a similar story to Prometheus - lean back and let **helm** do the work. As before, the main configuration is in **values.yaml**, where I spent the better part of a day making the data source and dashboard load on startup, so this service will work without additional configuration.

```
student@lab-kubernetes:~/scgc/grafana$ ./deploy-grafana.sh
namespace/monitoring unchanged
"bitnami" already exists with the same configuration, skipping
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "prometheus-community" chart repository
...Successfully got an update from the "bitnami" chart repository
Update Complete. #Happy Helming!#
NAME: grafana
LAST DEPLOYED: Fri May 12 14:01:02 2023
NAMESPACE: monitoring
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
CHART NAME: grafana
CHART VERSION: 8.4.2
APP VERSION: 9.5.2

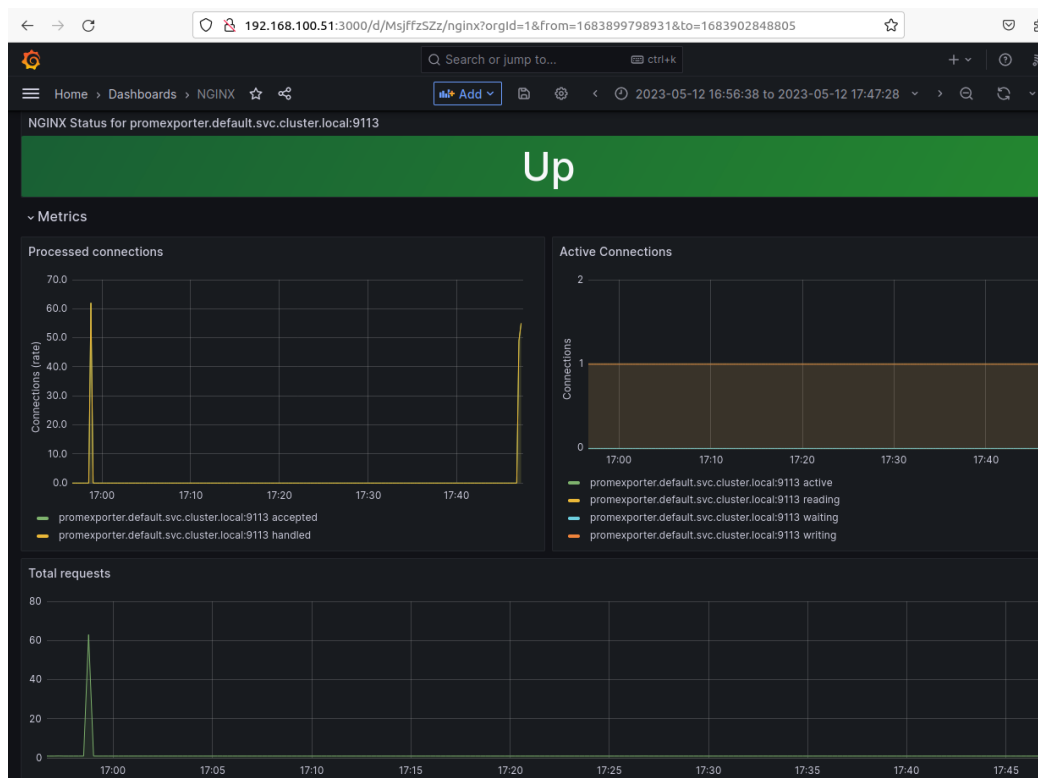
** Please be patient while the chart is being deployed **

1. Get the application URL by running these commands:
  echo "Browse to http://127.0.0.1:8080"
  kubectl port-forward svc/grafana 8080:3000 &

2. Get the admin credentials:

  echo "User: admin"
  echo "Password: $(kubectl get secret grafana-admin --namespace monitoring -o jsonpath="{.data.GF_SECURITY_ADMIN_PASSWORD}" | base64 -d)"
pod/grafana-998c48454-pn5h4 condition met
Credentials: admin/WGth4rSEEB
student@lab-kubernetes:~/scgc/grafana$ Forwarding from 0.0.0.0:3000 -> 3000
```

No better way to test an UI observability platform than by checking the UI after sending a bunch of requests to **nginx** like we did to test **task 1**.



6 Task 5 - Slowhttptest


This task serves to demonstrate how our deployment does against a Denial of Service attack. Keeping the server under a heavy load of 300 connections every second for 240 seconds led to the server being unresponsive for a few intervals.

Grafana reported the server down and no statistics were available in those intervals.



The downtime is better highlighted in the output file of `slowhttptest`.

Test parameters	
Test type	SLOW HEADERS
Number of connections	65539
Verb	GET
Content-Length header value	4096
Extra data max length	52
Interval between follow up data	1 seconds
Connections per seconds	300
Timeout for probe connection	3
Target test duration	240 seconds
Using proxy	no proxy




7 Task 6 - Secure Nginx

Knowing the attack model makes it easy to defend against it. Having built-in security features in both `nginx`¹ and `kubernetes` makes it even better.

The obvious ways to combat this were: caching `nginx` responses and scaling the service horizontally, both of which require extra resources. I opted instead for limiting the number of requests per client, limiting the number of connections per client, and closing slow connections, which do not bring additional costs and are well-suited for most applications.



Test parameters	
Test type	SLOW HEADERS
Number of connections	65539
Verb	GET
Content-Length header value	4096
Extra data max length	52
Interval between follow up data	1 seconds
Connections per seconds	300
Timeout for probe connection	3
Target test duration	240 seconds
Using proxy	no proxy



¹<https://www.nginx.com/blog/mitigating-ddos-attacks-with-nginx-and-nginx-plus/>

8 Mentions

To combat deployment failures caused by slow internet speed, deployment scripts download the required docker image and load it to the cluster.

Due to low space on the machine, I commented the majority of these additions, and only kept those that posed a problem on numerous redeploys e.g. `registry.k8s.io/kube-state-metrics/kube-state-metrics:v2.8.0`.